



OpenJDK / jdk8 / jdk8 / hotspot

view [src/share/vm/oops/markOop.hpp](#) @ 5820:87ee5ee27509

log
graph
tags
branches
changeset
browse

Added tag jdk8-b132 for changeset 0c94c41dcd70

author [katleman](#)

date Tue, 04 Mar 2014 11:51:03 -0800 (2014-03-05)

parents [da91efe96a93](#)

children

file

latest
diff
comparison
annotate
file log
raw
help

line source

line wrap: [on](#)

```

1 /*
2  * Copyright (c) 1997, 2012, Oracle and/or its affiliates. All rights reserved.
3  * DO NOT ALTER OR REMOVE COPYRIGHT NOTICES OR THIS FILE HEADER.
4  *
5  * This code is free software; you can redistribute it and/or modify it
6  * under the terms of the GNU General Public License version 2 only, as
7  * published by the Free Software Foundation.
8  *
9  * This code is distributed in the hope that it will be useful, but WITHOUT
10 * ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or
11 * FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License
12 * version 2 for more details (a copy is included in the LICENSE file that
13 * accompanied this code).
14 *
15 * You should have received a copy of the GNU General Public License version
16 * 2 along with this work; if not, write to the Free Software Foundation,
17 * Inc., 51 Franklin St, Fifth Floor, Boston, MA 02110-1301 USA.
18 *
19 * Please contact Oracle, 500 Oracle Parkway, Redwood Shores, CA 94065 USA
20 * or visit www.oracle.com if you need additional information or have any
21 * questions.
22 *
23 */
24
25 #ifndef SHARE_VM_OOPS_MARKOOP_HPP
26 #define SHARE_VM_OOPS_MARKOOP_HPP
27
28 #include "oops/oop.hpp"
```

```

29
30 // The markOop describes the header of an object.
31 //
32 // Note that the mark is not a real oop but just a word.
33 // It is placed in the oop hierarchy for historical reasons.
34 //
35 // Bit-format of an object header (most significant first, big endian layout below):
36 //
37 // 32 bits:
38 // -----
39 //          hash:25 ----->| age:4    biased_lock:1 lock:2 (normal object)
40 //          JavaThread*:23 epoch:2 age:4    biased_lock:1 lock:2 (biased object)
41 //          size:32 ----->| (CMS free block)
42 //          PromotedObject*:29 ----->| promo_bits:3 ----->| (CMS promoted object)
43 //
44 // 64 bits:
45 // -----
46 // unused:25 hash:31 -->| unused:1    age:4    biased_lock:1 lock:2 (normal object)
47 // JavaThread*:54 epoch:2 unused:1    age:4    biased_lock:1 lock:2 (biased object)
48 // PromotedObject*:61 ----->| promo_bits:3 ----->| (CMS promoted object)
49 // size:64 ----->| (CMS free block)
50 //
51 // unused:25 hash:31 -->| cms_free:1 age:4    biased_lock:1 lock:2 (C0OPs && normal object)
52 // JavaThread*:54 epoch:2 cms_free:1 age:4    biased_lock:1 lock:2 (C0OPs && biased object)
53 // narrowOop:32 unused:24 cms_free:1 unused:4 promo_bits:3 ----->| (C0OPs && CMS promoted object)
54 // unused:21 size:35 -->| cms_free:1 unused:7 ----->| (C0OPs && CMS free block)
55 //
56 // - hash contains the identity hash value: largest value is
57 //   31 bits, see os::random(). Also, 64-bit vm's require
58 //   a hash value no bigger than 32 bits because they will not
59 //   properly generate a mask larger than that: see library_call.cpp
60 //   and cl_CodePatterns_sparc.cpp.
61 //
62 // - the biased lock pattern is used to bias a lock toward a given
63 //   thread. When this pattern is set in the low three bits, the lock
64 //   is either biased toward a given thread or "anonymously" biased,
65 //   indicating that it is possible for it to be biased. When the
66 //   lock is biased toward a given thread, locking and unlocking can
67 //   be performed by that thread without using atomic operations.
68 //   When a lock's bias is revoked, it reverts back to the normal
69 //   locking scheme described below.

```

```

70 //
71 //     Note that we are overloading the meaning of the "unlocked" state
72 //     of the header. Because we steal a bit from the age we can
73 //     guarantee that the bias pattern will never be seen for a truly
74 //     unlocked object.
75 //
76 //     Note also that the biased state contains the age bits normally
77 //     contained in the object header. Large increases in scavenge
78 //     times were seen when these bits were absent and an arbitrary age
79 //     assigned to all biased objects, because they tended to consume a
80 //     significant fraction of the eden semispaces and were not
81 //     promoted promptly, causing an increase in the amount of copying
82 //     performed. The runtime system aligns all JavaThread* pointers to
83 //     a very large value (currently 128 bytes (32bVM) or 256 bytes (64bVM))
84 //     to make room for the age bits & the epoch bits (used in support of
85 //     biased locking), and for the CMS "freeness" bit in the 64bVM (+C00Ps).
86 //
87 //     [JavaThread* | epoch | age | 1 | 01]      lock is biased toward given thread
88 //     [0           | epoch | age | 1 | 01]      lock is anonymously biased
89 //
90 // - the two lock bits are used to describe three states: locked/unlocked and monitor.
91 //
92 //     [ptr          | 00] locked          ptr points to real header on stack
93 //     [header       | 0 | 01] unlocked      regular object header
94 //     [ptr          | 10] monitor          inflated lock (header is wapped out)
95 //     [ptr          | 11] marked           used by markSweep to mark an object
96 //                                           not valid at any other time
97 //
98 //     We assume that stack/thread pointers have the lowest two bits cleared.
99
100 class BasicLock;
101 class ObjectMonitor;
102 class JavaThread;
103
104 class markOopDesc: public oopDesc {
105 private:
106     // Conversion
107     uintptr_t value() const { return (uintptr_t) this; }
108
109 public:
110     // Constants

```

```

111 enum { age_bits          = 4,
112         lock_bits        = 2,
113         biased_lock_bits = 1,
114         max_hash_bits    = BitsPerWord - age_bits - lock_bits - biased_lock_bits,
115         hash_bits        = max_hash_bits > 31 ? 31 : max_hash_bits,
116         cms_bits         = LP64_ONLY(1) NOT_LP64(0),
117         epoch_bits       = 2
118 };
119
120 // The biased locking code currently requires that the age bits be
121 // contiguous to the lock bits.
122 enum { lock_shift        = 0,
123         biased_lock_shift = lock_bits,
124         age_shift         = lock_bits + biased_lock_bits,
125         cms_shift         = age_shift + age_bits,
126         hash_shift        = cms_shift + cms_bits,
127         epoch_shift       = hash_shift
128 };
129
130 enum { lock_mask          = right_n_bits(lock_bits),
131         lock_mask_in_place = lock_mask << lock_shift,
132         biased_lock_mask   = right_n_bits(lock_bits + biased_lock_bits),
133         biased_lock_mask_in_place = biased_lock_mask << lock_shift,
134         biased_lock_bit_in_place = 1 << biased_lock_shift,
135         age_mask           = right_n_bits(age_bits),
136         age_mask_in_place  = age_mask << age_shift,
137         epoch_mask         = right_n_bits(epoch_bits),
138         epoch_mask_in_place = epoch_mask << epoch_shift,
139         cms_mask           = right_n_bits(cms_bits),
140         cms_mask_in_place  = cms_mask << cms_shift
141 #ifndef _WIN64
142         ,hash_mask         = right_n_bits(hash_bits),
143         hash_mask_in_place = (address_word)hash_mask << hash_shift
144 #endif
145 };
146
147 // Alignment of JavaThread pointers encoded in object header required by biased locking
148 enum { biased_lock_alignment = 2 << (epoch_shift + epoch_bits)
149 };
150
151 #ifdef _WIN64

```

```
152 // These values are too big for Win64
153 const static uintptr_t hash_mask = right_n_bits(hash_bits);
154 const static uintptr_t hash_mask_in_place =
155     (address_word)hash_mask << hash_shift;
156 #endif
157
158 enum { locked_value          = 0,
159        unlocked_value        = 1,
160        monitor_value         = 2,
161        marked_value          = 3,
162        biased_lock_pattern    = 5
163 };
164
165 enum { no_hash                = 0 }; // no hash value assigned
166
167 enum { no_hash_in_place       = (address_word)no_hash << hash_shift,
168        no_lock_in_place       = unlocked_value
169 };
170
171 enum { max_age                = age_mask };
172
173 enum { max_bias_epoch         = epoch_mask };
174
175 // Biased Locking accessors.
176 // These must be checked by all code which calls into the
177 // ObjectSynchronizer and other code. The biasing is not understood
178 // by the lower-level CAS-based locking code, although the runtime
179 // fixes up biased locks to be compatible with it when a bias is
180 // revoked.
181 bool has_bias_pattern() const {
182     return (mask_bits(value(), biased_lock_mask_in_place) == biased_lock_pattern);
183 }
184 JavaThread* biased_locker() const {
185     assert(has_bias_pattern(), "should not call this otherwise");
186     return (JavaThread*) ((intptr_t) (mask_bits(value(), ~(biased_lock_mask_in_place | age_mask_in_place |
epoch_mask_in_place))));
187 }
188 // Indicates that the mark has the bias bit set but that it has not
189 // yet been biased toward a particular thread
190 bool is_biased_anonymously() const {
```

```
191     return (has_bias_pattern() && (biased_locker() == NULL));
192 }
193 // Indicates epoch in which this bias was acquired. If the epoch
194 // changes due to too many bias revocations occurring, the biases
195 // from the previous epochs are all considered invalid.
196 int bias_epoch() const {
197     assert(has_bias_pattern(), "should not call this otherwise");
198     return (mask_bits(value(), epoch_mask_in_place) >> epoch_shift);
199 }
200 markOop set_bias_epoch(int epoch) {
201     assert(has_bias_pattern(), "should not call this otherwise");
202     assert((epoch & (~epoch_mask)) == 0, "epoch overflow");
203     return markOop(mask_bits(value(), ~epoch_mask_in_place) | (epoch << epoch_shift));
204 }
205 markOop incr_bias_epoch() {
206     return set_bias_epoch((1 + bias_epoch()) & epoch_mask);
207 }
208 // Prototype mark for initialization
209 static markOop biased_locking_prototype() {
210     return markOop( biased_lock_pattern );
211 }
212
213 // lock accessors (note that these assume lock_shift == 0)
214 bool is_locked() const {
215     return (mask_bits(value(), lock_mask_in_place) != unlocked_value);
216 }
217 bool is_unlocked() const {
218     return (mask_bits(value(), biased_lock_mask_in_place) == unlocked_value);
219 }
220 bool is_marked() const {
221     return (mask_bits(value(), lock_mask_in_place) == marked_value);
222 }
223 bool is_neutral() const { return (mask_bits(value(), biased_lock_mask_in_place) == unlocked_value); }
224
225 // Special temporary state of the markOop while being inflated.
226 // Code that looks at mark outside a lock need to take this into account.
227 bool is_being_inflated() const { return (value() == 0); }
228
229 // Distinguished markword value - used when inflating over
230 // an existing stacklock. 0 indicates the markword is "BUSY".
231 // Lockword mutators that use a LD...CAS idiom should always
```

```
232 // check for and avoid overwriting a 0 value installed by some
233 // other thread. (They should spin or block instead. The 0 value
234 // is transient and *should* be short-lived).
235 static markOop INFLATING() { return (markOop) 0; } // inflate-in-progress
236
237 // Should this header be preserved during GC?
238 inline bool must_be_preserved(oop obj_containing_mark) const;
239 inline bool must_be_preserved_with_bias(oop obj_containing_mark) const;
240
241 // Should this header (including its age bits) be preserved in the
242 // case of a promotion failure during scavenge?
243 // Note that we special case this situation. We want to avoid
244 // calling BiasedLocking::preserve_marks()/restore_marks() (which
245 // decrease the number of mark words that need to be preserved
246 // during GC) during each scavenge. During scavenges in which there
247 // is no promotion failure, we actually don't need to call the above
248 // routines at all, since we don't mutate and re-initialize the
249 // marks of promoted objects using init_mark(). However, during
250 // scavenges which result in promotion failure, we do re-initialize
251 // the mark words of objects, meaning that we should have called
252 // these mark word preservation routines. Currently there's no good
253 // place in which to call them in any of the scavengers (although
254 // guarded by appropriate locks we could make one), but the
255 // observation is that promotion failures are quite rare and
256 // reducing the number of mark words preserved during them isn't a
257 // high priority.
258 inline bool must_be_preserved_for_promotion_failure(oop obj_containing_mark) const;
259 inline bool must_be_preserved_with_bias_for_promotion_failure(oop obj_containing_mark) const;
260
261 // Should this header be preserved during a scavenge where CMS is
262 // the old generation?
263 // (This is basically the same body as must_be_preserved_for_promotion_failure(),
264 // but takes the Klass* as argument instead)
265 inline bool must_be_preserved_for_cms_scavenge(Klass* klass_of_obj_containing_mark) const;
266 inline bool must_be_preserved_with_bias_for_cms_scavenge(Klass* klass_of_obj_containing_mark) const;
267
268 // WARNING: The following routines are used EXCLUSIVELY by
269 // synchronization functions. They are not really gc safe.
270 // They must get updated if markOop layout get changed.
271 markOop set_unlocked() const {
272     return markOop(value() | unlocked_value);
```

```
273 }
274 bool has_locker() const {
275     return ((value() & lock_mask_in_place) == locked_value);
276 }
277 BasicLock* locker() const {
278     assert(has_locker(), "check");
279     return (BasicLock*) value();
280 }
281 bool has_monitor() const {
282     return ((value() & monitor_value) != 0);
283 }
284 ObjectMonitor* monitor() const {
285     assert(has_monitor(), "check");
286     // Use xor instead of &~ to provide one extra tag-bit check.
287     return (ObjectMonitor*) (value() ^ monitor_value);
288 }
289 bool has_displaced_mark_helper() const {
290     return ((value() & unlocked_value) == 0);
291 }
292 markOop displaced_mark_helper() const {
293     assert(has_displaced_mark_helper(), "check");
294     intptr_t ptr = (value() & ~monitor_value);
295     return *(markOop*)ptr;
296 }
297 void set_displaced_mark_helper(markOop m) const {
298     assert(has_displaced_mark_helper(), "check");
299     intptr_t ptr = (value() & ~monitor_value);
300     *(markOop*)ptr = m;
301 }
302 markOop copy_set_hash(intptr_t hash) const {
303     intptr_t tmp = value() & (~hash_mask_in_place);
304     tmp |= ((hash & hash_mask) << hash_shift);
305     return (markOop)tmp;
306 }
307 // it is only used to be stored into BasicLock as the
308 // indicator that the lock is using heavyweight monitor
309 static markOop unused_mark() {
310     return (markOop) marked_value;
311 }
312 // the following two functions create the markOop to be
313 // stored into object header, it encodes monitor info
314 static markOop encode(BasicLock* lock) {
```



```
315     return (markOop) lock;
316 }
317 static markOop encode(ObjectMonitor* monitor) {
318     intptr_t tmp = (intptr_t) monitor;
319     return (markOop) (tmp | monitor_value);
320 }
321 static markOop encode(JavaThread* thread, uint age, int bias_epoch) {
322     intptr_t tmp = (intptr_t) thread;
323     assert(UseBiasedLocking && ((tmp & (epoch_mask_in_place | age_mask_in_place | biased_lock_mask_in_place)) == 0),
"misaligned JavaThread pointer");
324     assert(age <= max_age, "age too large");
325     assert(bias_epoch <= max_bias_epoch, "bias epoch too large");
326     return (markOop) (tmp | (bias_epoch << epoch_shift) | (age << age_shift) | biased_lock_pattern);
327 }
328
329 // used to encode pointers during GC
330 markOop clear_lock_bits() { return markOop(value() & ~lock_mask_in_place); }
331
332 // age operations
333 markOop set_marked() { return markOop((value() & ~lock_mask_in_place) | marked_value); }
334 markOop set_unmarked() { return markOop((value() & ~lock_mask_in_place) | unlocked_value); }
335
336 uint age() const { return mask_bits(value() >> age_shift, age_mask); }
337 markOop set_age(uint v) const {
338     assert((v & ~age_mask) == 0, "shouldn't overflow age field");
339     return markOop((value() & ~age_mask_in_place) | (((uintptr_t)v & age_mask) << age_shift));
340 }
341 markOop incr_age() const { return age() == max_age ? markOop(this) : set_age(age() + 1); }
342
343 // hash operations
344 intptr_t hash() const {
345     return mask_bits(value() >> hash_shift, hash_mask);
346 }
347
348 bool has_no_hash() const {
349     return hash() == no_hash;
350 }
351
352 // Prototype mark for initialization
353 static markOop prototype() {
```

```
354     return markOop( no_hash_in_place | no_lock_in_place );
355 }
356
357 // Helper function for restoration of unmarked mark oops during GC
358 static inline markOop prototype_for_object(oop obj);
359
360 // Debugging
361 void print_on(outputStream* st) const;
362
363 // Prepare address of oop for placement into mark
364 inline static markOop encode_pointer_as_mark(void* p) { return markOop(p)->set_marked(); }
365
366 // Recover address of oop from encoded form used in mark
367 inline void* decode_pointer() { if (UseBiasedLocking && has_bias_pattern()) return NULL; return clear_lock_bits(); }
368
369 // These markOops indicate cms free chunk blocks and not objects.
370 // In 64 bit, the markOop is set to distinguish them from oops.
371 // These are defined in 32 bit mode for vmStructs.
372 const static uintptr_t cms_free_chunk_pattern = 0x1;
373
374 // Constants for the size field.
375 enum { size_shift          = cms_shift + cms_bits,
376        size_bits          = 35      // need for compressed oops 32G
377        };
378 // These values are too big for Win64
379 const static uintptr_t size_mask = LP64_ONLY(right_n_bits(size_bits))
380                                     NOT_LP64(0);
381 const static uintptr_t size_mask_in_place =
382                                     (address_word)size_mask << size_shift;
383
384 #ifdef _LP64
385 static markOop cms_free_prototype() {
386     return markOop(((intptr_t)prototype() & ~cms_mask_in_place) |
387                    ((cms_free_chunk_pattern & cms_mask) << cms_shift));
388 }
389 uintptr_t cms_encoding() const {
390     return mask_bits(value() >> cms_shift, cms_mask);
391 }
392 bool is_cms_free_chunk() const {
393     return is_neutral() &&
```

```
394         (cms_encoding() & cms_free_chunk_pattern) == cms_free_chunk_pattern;
395     }
396
397     size_t get_size() const          { return (size_t)(value() >> size_shift); }
398     static markOop set_size_and_free(size_t size) {
399         assert((size & ~size_mask) == 0, "shouldn't overflow size field");
400         return markOop(((intptr_t)cms_free_prototype() & ~size_mask_in_place) |
401                        (((intptr_t)size & size_mask) << size_shift));
402     }
403 #endif // _LP64
404 };
405
406 #endif // SHARE_VM_OOPS_MARKOOP_HPP
```

© 2007, 2020 Oracle and/or its affiliates
Terms of Use · Privacy · Trademarks